

CS 111: Homework 4: Due by 11:59 pm Sunday, October 24, 2021

Submit your paper as one PDF file, and tell GradeScope which page(s) each problem is on. If you worked with a partner, you must each turn in your own homework paper, and report the name and perm number of your partner. No groups of more than two allowed.

1. Do problem 2.3 on pages 32–33 of the NCM book, showing the `numpy` code you use and its output. Note: To understand intuitively what the problem means by “assume that joint 1 is rigidly fixed both horizontally and vertically and that joint 8 is fixed vertically,” think of the truss as a (2-dimensional) drawbridge across a river, with the left end being a hinge and the right end lying on the ground.

2. Recall that a symmetric matrix A is *positive definite* (SPD for short) if and only if $x^T A x > 0$ for every nonzero vector x .

2.1 Find a 2-by-2 matrix A that (1) is symmetric, (2) is not singular, and (3) has all its elements greater than zero, but (4) is *not* SPD. Show a nonzero vector x such that $x^T A x < 0$.

2.2 Let B be an m -by- n matrix (m and n may or may not be equal) whose rank is n . Prove that the matrix $A = B^T B$ is SPD (mathematically, not experimentally).

3. If A is symmetric, we don’t need to store all n^2 of its elements; we can just store the $n(n+1)/2$ elements of the upper triangle of A , for example. If A is symmetric and also positive definite then there is a symmetric version of Gaussian elimination called *Cholesky factorization*. You can read about Cholesky and his factorization in NCM problem 2.5 (pages 35–36), but don’t do that problem.

The Cholesky factorization of an SPD matrix is

$$A = R^T R,$$

where R is an upper triangular matrix with all its diagonal elements positive. Notice that there’s only one triangular matrix R involved, so computing the factorization should only need to compute $n(n+1)/2$ numbers, not n^2 numbers like LU factorization.

One way to get R from A is to factor $A = LU$ with no pivoting (there’s a theorem that says this is possible, and stable, if A is SPD); then write $U = DV$ where D is diagonal and V is upper triangular with ones on the diagonal; then show that $L = V^T$ so that $A = V^T D V$; then finally take $R = \sqrt{D} V$, where \sqrt{D} is just the diagonal matrix of square roots of diagonal elements of D ; then we have $A = V^T D V = R^T R$ as desired. However, this method does twice as much work as it needs to, because it computes all n^2 elements of L and U .

Your assignment is to write a routine `R = Cfactor(A)` that returns the factor R without ever touching the lower triangle of A or the lower triangle of R (or any other n -by- n matrix). For full credit, your routine should also only do about half as many arithmetic operations as `L, U = cs111.LUfactorNoPiv(A)`. For debugging, you can generate a random n -by- n SPD matrix A by saying

```
B = np.random.randn(n, n)
A = B.T @ B
```

Explain in English (in LaTeX) how your `Cfactor()` works. Demonstrate that it works by generating a 10-by-10 SPD matrix A as above, generating a random 10-vector b , and comparing the solution to $Ax = b$ from `x = cs111.LUsolve(A,b)` to the solution you get by saying

```
R = Cfactor(A)
y = cs111.Lsolve(R.T, b)
x = cs111.Lsolve(R, b)
```

Finally, do an experiment to compare the running time of your `Cfactor(A)` with that of `LUfactorNoPiv(A)`, for a range of values of n up to large enough that the routines take a few seconds to run. Report your running times, and make a plot of the ratio of `Cfactor(A)` time to `LUfactorNoPiv(A)` time against n . (You can time one line of code in Jupyter by saying `%time line-of-code`, or you can time a whole window by starting it with `%%time`.)

4. Here you will experiment with solving $Ax = b$ using various solvers from class and from `numpy`. For this problem, you should use the 3-D version of the temperature matrix from `make_A_3D()`. You can use the version of `make_A_3D()` you wrote for Homework 3, or if you prefer you can use my version (which is in the latest update of `cs111/temperature.py` on Gauchospace). For a right-hand side b , use the vector of row sums, `b = A @ np.ones(n)`, so that you know that the exact solution to $Ax = b$ is the vector of all ones.

Experiment with solving $Ax = b$ for the temperature x , for various values of k , using five different solvers as follows. For each solver, you should report (showing code and output) the largest value of k for which that solver could solve $Ax = b$ within 30 seconds. For all but the last solver, use the sparse version of A from `make_A_3D()`.

- The `cs111.CGsolve()` conjugate gradient solver, from class. (You can vary the arguments `tol` and `max_iters` to make it find a more accurate solution.)
- The `cs111.Jsolve()` Jacobi solver, also from class. (Again you can vary `tol` and `max_iters`.)
- The `scipy` sparse conjugate gradient solver `scipy.sparse.linalg.cg()`.
- The `scipy` sparse LU solver `scipy.sparse.linalg.spsolve()`.
- The dense LU solver `cs111.LUsolve()` from class. For this solver, you will have to convert A to a dense array with `A.toarray()`. Warning! This will run out of memory if k gets very big.

For each solve, measure and report the run time, the relative residual norm, and the relative error norm $\|x_{\text{exact}} - x\|/\|x_{\text{exact}}\|$. Which solvers are more accurate? Which are faster? How do the answers to these questions change as you change k ?

Warning: Start with very small values of k , and be cautious as you increase k ! The matrices get big in a hurry. Different solvers will fall over for different values of k .